

Assignment 5

The goal of this assignment is to train a Word2Vec skip-gram model over [Text8](#) data.

In [1]:

*# These are all the modules we'll be using later. Make sure you can import them
before proceeding further.*

```
%matplotlib inline
from __future__ import print_function
import collections
import math
import numpy as np
import os
import random
import tensorflow as tf
import zipfile
from matplotlib import pylab
from six.moves import range
from six.moves.urllib.request import urlretrieve
from sklearn.manifold import TSNE
```

Download the data from the source website if necessary.

In [2]:

```
url = 'http://mattmahoney.net/dc/'

def maybe_download(filename, expected_bytes):
    """Download a file if not present, and make sure it's the right size."""
    if not os.path.exists(filename):
        filename, _ = urlretrieve(url + filename, filename)
    statinfo = os.stat(filename)
    if statinfo.st_size == expected_bytes:
        print('Found and verified %s' % filename)
    else:
        print(statinfo.st_size)
        raise Exception(
            'Failed to verify ' + filename + '. Can you get to it with a browser?')
    return filename
```

```
filename = maybe_download('text8.zip', 31344016)
```

Found and verified text8.zip

Read the data into a string.

In [3]:

```
def read_data(filename):
    """Extract the first file enclosed in a zip file as a list of words"""
    with zipfile.ZipFile(filename) as f:
        data = tf.compat.as_str(f.read(f.namelist()[0])).split()
    return data
```

```
words = read_data(filename)
print('Data size %d' % len(words))
```

Data size 17005207

Build the dictionary and replace rare words with UNK token.

In [4]:

```
vocabulary_size = 50000
```

```
def build_dataset(words):
    count = [['UNK', -1]]
    count.extend(collections.Counter(words).most_common(vocabulary_size - 1))
    dictionary = dict()
    for word, _ in count:
        dictionary[word] = len(dictionary)
    data = list()
    unk_count = 0
    for word in words:
        if word in dictionary:
            index = dictionary[word]
        else:
            index = 0 # dictionary['UNK']
            unk_count = unk_count + 1
        data.append(index)
    count[0][1] = unk_count
    reverse_dictionary = dict(zip(dictionary.values(), dictionary.keys()))
    return data, count, dictionary, reverse_dictionary
```

```
data, count, dictionary, reverse_dictionary = build_dataset(words)
```

```
print('Most common words (+UNK)', count[:5])
```

```
print('Sample data', data[:10])
```

```
del words # Hint to reduce memory.
```

```
Most common words (+UNK) [['UNK', 418391], ('the', 1061396), ('of', 593677),
('and', 416629), ('one', 411764)]
```

```
Sample data [5239, 3084, 12, 6, 195, 2, 3137, 46, 59, 156]
```

Function to generate a training batch for the skip-gram model.

In [5]:

```
data_index = 0
```

```
def generate_batch(batch_size, num_skips, skip_window):
    global data_index
    assert batch_size % num_skips == 0
    assert num_skips <= 2 * skip_window
    batch = np.ndarray(shape=(batch_size), dtype=np.int32)
    labels = np.ndarray(shape=(batch_size, 1), dtype=np.int32)
    span = 2 * skip_window + 1 # [ skip_window target skip_window ]
    buffer = collections.deque(maxlen=span)

    for _ in range(span):
        buffer.append(data[data_index])
        data_index = (data_index + 1) % len(data)

    for i in range(batch_size // num_skips):
        target = skip_window # target label at the center of the buffer
        targets_to_avoid = [ skip_window ]
```

```

    for j in range(num_skips):
        while target in targets_to_avoid:
            target = random.randint(0, span - 1)
        targets_to_avoid.append(target)
        batch[i * num_skips + j] = buffer[skip_window]
        labels[i * num_skips + j, 0] = buffer[target]
        buffer.append(data[data_index])
        data_index = (data_index + 1) % len(data)
    return batch, labels

print('data:', [reverse_dictionary[di] for di in data[:8]])

for num_skips, skip_window in [(2, 1), (4, 2)]:
    data_index = 0
    batch, labels = generate_batch(batch_size=8, num_skips=num_skips,
    skip_window=skip_window)
    print('\nwith num_skips = %d and skip_window = %d:' % (num_skips,
    skip_window))
    print('    batch:', [reverse_dictionary[bi] for bi in batch])
    print('    labels:', [reverse_dictionary[li] for li in labels.reshape(8)])
data: ['anarchism', 'originated', 'as', 'a', 'term', 'of', 'abuse', 'first']

with num_skips = 2 and skip_window = 1:
    batch: ['originated', 'originated', 'as', 'as', 'a', 'a', 'term', 'term']
    labels: ['as', 'anarchism', 'originated', 'a', 'as', 'term', 'of', 'a']

with num_skips = 4 and skip_window = 2:
    batch: ['as', 'as', 'as', 'as', 'a', 'a', 'a', 'a']
    labels: ['term', 'a', 'originated', 'anarchism', 'term', 'of',
    'originated', 'as']

```

Train a skip-gram model.

In [6]:

```

batch_size = 128
embedding_size = 128 # Dimension of the embedding vector.
skip_window = 1 # How many words to consider left and right.
num_skips = 2 # How many times to reuse an input to generate a label.
# We pick a random validation set to sample nearest neighbors. Here we limit the
# validation samples to the words that have a low numeric ID, which by
# construction are also the most frequent.
valid_size = 16 # Random set of words to evaluate similarity on.
valid_window = 100 # Only pick dev samples in the head of the distribution.
valid_examples = np.array(random.sample(range(valid_window), valid_size))
num_sampled = 64 # Number of negative examples to sample.
graph = tf.Graph()

with graph.as_default(), tf.device('/cpu:0'):

    # Input data.
    train_dataset = tf.placeholder(tf.int32, shape=[batch_size])
    train_labels = tf.placeholder(tf.int32, shape=[batch_size, 1])
    valid_dataset = tf.constant(valid_examples, dtype=tf.int32)

```

```

# Variables.
embeddings = tf.Variable(
    tf.random_uniform([vocabulary_size, embedding_size], -1.0, 1.0))
softmax_weights = tf.Variable(
    tf.truncated_normal([vocabulary_size, embedding_size],
                        stddev=1.0 / math.sqrt(embedding_size)))
softmax_biases = tf.Variable(tf.zeros([vocabulary_size]))

# Model.
# Look up embeddings for inputs.
embed = tf.nn.embedding_lookup(embeddings, train_dataset)
# Compute the softmax loss, using a sample of the negative labels each time.
loss = tf.reduce_mean(
    tf.nn.sampled_softmax_loss(softmax_weights, softmax_biases, embed,
                                train_labels, num_sampled, vocabulary_size))

# Optimizer.
# Note: The optimizer will optimize the softmax_weights AND the embeddings.
# This is because the embeddings are defined as a variable quantity and the
# optimizer's `minimize` method will by default modify all variable
quantities
# that contribute to the tensor it is passed.
# See docs on `tf.train.Optimizer.minimize()` for more details.
optimizer = tf.train.AdagradOptimizer(1.0).minimize(loss)

# Compute the similarity between minibatch examples and all embeddings.
# We use the cosine distance:
norm = tf.sqrt(tf.reduce_sum(tf.square(embeddings), 1, keep_dims=True))
normalized_embeddings = embeddings / norm
valid_embeddings = tf.nn.embedding_lookup(
    normalized_embeddings, valid_dataset)
similarity = tf.matmul(valid_embeddings, tf.transpose(normalized_embeddings))
                                                                    In [7]:

num_steps = 100001

with tf.Session(graph=graph) as session:
    tf.initialize_all_variables().run()
    print('Initialized')
    average_loss = 0

    for step in range(num_steps):
        batch_data, batch_labels = generate_batch(
            batch_size, num_skips, skip_window)
        feed_dict = {train_dataset : batch_data, train_labels : batch_labels}
        _, l = session.run([optimizer, loss], feed_dict=feed_dict)
        average_loss += l

    if step % 2000 == 0:
        if step > 0:
            average_loss = average_loss / 2000
            # The average loss is an estimate of the loss over the last 2000 batches.

```

```

print('Average loss at step %d: %f' % (step, average_loss))
average_loss = 0
# note that this is expensive (~20% slowdown if computed every 500 steps)

if step % 10000 == 0:
    sim = similarity.eval()

    for i in range(valid_size):
        valid_word = reverse_dictionary[valid_examples[i]]
        top_k = 8 # number of nearest neighbors
        nearest = (-sim[i, :]).argsort()[1:top_k+1]
        log = 'Nearest to %s:' % valid_word

        for k in range(top_k):
            close_word = reverse_dictionary[nearest[k]]
            log = '%s %s,' % (log, close_word)
        print(log)
    final_embeddings = normalized_embeddings.eval()

Initialized
Average loss at step 0: 8.161867
Nearest to all: ajaccio, legged, creationism, vagaries, guitarists, penetration,
ibero, mactutor,
Nearest to history: diaeresis, nicolai, canmore, mudslides, harmonize, excerpt,
thujone, fried,
Nearest to from: eclac, texan, subsisting, kramer, arminians, traction,
displayed, dissatisfied,
Nearest to system: arthur, gmt, midair, zug, thicker, pathological, fabricated,
kabuki,
Nearest to or: omnium, condemns, lama, keen, capitulated, participate, removes,
aglaulus,
.....
Nearest to often: sometimes, usually, commonly, generally, frequently, now,
Nearest to four: six, seven, eight, five, three, nine, two, zero,
Nearest to new: different, painterly, huge, masoretic, mohiam, scuba, outback,
old,
Nearest to there: they, it, he, we, still, often, usually, she,

```

In [8]:

```

num_points = 400

tsne = TSNE(perplexity=30, n_components=2, init='pca', n_iter=5000)
two_d_embeddings = tsne.fit_transform(final_embeddings[1:num_points+1, :])

```

In [9]:

```

def plot(embeddings, labels):
    assert embeddings.shape[0] >= len(labels), 'More labels than embeddings'
    pylab.figure(figsize=(15,15)) # in inches
    for i, label in enumerate(labels):
        x, y = embeddings[i,:]
        pylab.scatter(x, y)
        pylab.annotate(label, xy=(x, y), xytext=(5, 2), textcoords='offset points',
                        ha='right', va='bottom')
    pylab.show()

```

```
if self. edgecolors == str('face'):
```



In [9]:

```
%matplotlib inline
from __future__ import print_function
import collections
import math
import numpy as np
import os
import random
import tensorflow as tf
import zipfile
from matplotlib import pylab
from six.moves import range
from six.moves.urllib.request import urlretrieve
from sklearn.manifold import TSNE
from itertools import compress
```

```
#data preparation
```

```
url = 'http://matmahoney.net/dc/'
```

```
def maybe_download(filename, expected_bytes):
    """Download a file if not present, and make sure it's the right size."""
    if not os.path.exists(filename):
        filename, _ = urlretrieve(url + filename, filename)
    statinfo = os.stat(filename)
    if statinfo.st_size == expected_bytes:
        print('Found and verified %s' % filename)
    else:
        print(statinfo.st_size)
        raise Exception(
            'Failed to verify ' + filename + '. Can you get to it with a browser?')
    return filename
```

```
filename = maybe_download('text8.zip', 31344016)
```

```
def read_data(filename):
    """Extract the first file enclosed in a zip file as a list of words"""
    with zipfile.ZipFile(filename) as f:
        data = tf.compat.as_str(f.read(f.namelist()[0])).split()
    return data
```

```
words = read_data(filename)
print('Data size %d' % len(words))
```

```
vocabulary_size = 50000
```

```
def build_dataset(words):
    count = [['UNK', -1]]
    count.extend(collections.Counter(words).most_common(vocabulary_size - 1))
    dictionary = dict()
    for word, _ in count:
        dictionary[word] = len(dictionary)
    data = list()
    unk_count = 0
    for word in words:
        if word in dictionary:
            index = dictionary[word]
        else:
            index = 0 # dictionary['UNK']
            unk_count = unk_count + 1
        data.append(index)
    count[0][1] = unk_count
    reverse_dictionary = dict(zip(dictionary.values(), dictionary.keys()))
    return data, count, dictionary, reverse_dictionary
```

```
data, count, dictionary, reverse_dictionary = build_dataset(words)
print('Most common words (+UNK)', count[:5])
```

```
print('Sample data', data[:10])
del words # Hint to reduce memory.
```

Found and verified text8.zip

Data size 17005207

Most common words (+UNK) [['UNK', 418391], ('the', 1061396), ('of', 593677), ('and', 416629), ('one', 411764)]

Sample data [5239, 3084, 12, 6, 195, 2, 3137, 46, 59, 156]

In [20]:

```
data_index = 0
```

#generally what we do here is build data as a pack of words to left and right of the target word

#and label is the target word

```
def generate_batch_cbow(batch_size, bag_window):
    global data_index
    span = 2 * bag_window + 1 # [ bag_window target bag_window ]
    batch = np.ndarray(shape=(batch_size, span - 1), dtype=np.int32)
    labels = np.ndarray(shape=(batch_size, 1), dtype=np.int32)
    buffer = collections.deque(maxlen=span)

    for _ in range(span):
        buffer.append(data[data_index])
        data_index = (data_index + 1) % len(data)

    for i in range(batch_size):
        # just for testing
        buffer_list = list(buffer)
        labels[i, 0] = buffer_list.pop(bag_window)
        batch[i] = buffer_list
        # iterate to the next buffer
        buffer.append(data[data_index])
        data_index = (data_index + 1) % len(data)

    return batch, labels
```

```
print('data:', [reverse_dictionary[di] for di in data[:16]])
```

```
for bag_window in [1, 2]:
    data_index = 0
    batch, labels = generate_batch_cbow(batch_size=4, bag_window=bag_window)
    print('\nwith bag_window = %d:' % (bag_window))
    print('    batch:', [[reverse_dictionary[w] for w in bi] for bi in batch])
    print('    labels:', [reverse_dictionary[li] for li in labels.reshape(4)])
```

```
data: ['anarchism', 'originated', 'as', 'a', 'term', 'of', 'abuse', 'first',
'used', 'against', 'early', 'working', 'class', 'radicals', 'including', 'the']
```

```
with bag_window = 1:
```

```
batch: [['anarchism', 'as'], ['originated', 'a'], ['as', 'term'], ['a', 'of']]
labels: ['originated', 'as', 'a', 'term']
```



```
with bag_window = 2:  
    batch: [['anarchism', 'originated', 'a', 'term'], ['originated', 'as', 'term',  
'of'], ['as', 'a', 'of', 'abuse'], ['a', 'term', 'abuse', 'first']]  
    labels: ['as', 'a', 'term', 'of']
```

In [21]:

```
batch_size = 128  
embedding_size = 128 # Dimension of the embedding vector.  
bag_window = 2 # How many words to consider left and right.  
# We pick a random validation set to sample nearest neighbors. here we limit  
the  
# validation samples to the words that have a low numeric ID, which by  
# construction are also the most frequent.  
valid_size = 16 # Random set of words to evaluate similarity on.  
valid_window = 100 # Only pick dev samples in the head of the distribution.  
valid_examples = np.array(random.sample(range(valid_window), valid_size))  
num_sampled = 64 # Number of negative examples to sample.
```

```
graph = tf.Graph()
```

```
with graph.as_default(), tf.device('/cpu:0'):
```

```
    # Input data.
```

```
    train_dataset = tf.placeholder(tf.int32, shape=[batch_size, bag_window*2])  
    train_labels = tf.placeholder(tf.int32, shape=[batch_size, 1])  
    valid_dataset = tf.constant(valid_examples, dtype=tf.int32)
```

```
    # Variables.
```

```
    embeddings = tf.Variable(  
        tf.random_uniform([vocabulary_size, embedding_size], -1.0, 1.0))  
    softmax_weights = tf.Variable(  
        tf.truncated_normal([vocabulary_size, embedding_size],  
                             stddev=1.0 / math.sqrt(embedding_size)))  
    softmax_biases = tf.Variable(tf.zeros([vocabulary_size]))
```

```
    # Model.
```

```
    # Look up embeddings for inputs.
```

```
    embeds = tf.nn.embedding_lookup(embeddings, train_dataset)
```

```
    # Compute the softmax loss, using a sample of the negative labels each time.
```

```
    # Notice per description of CBOW we have to sum all the embeddings!
```

```
    loss = tf.reduce_mean(  
        tf.nn.sampled_softmax_loss(softmax_weights, softmax_biases,  
        tf.reduce_sum(embeds,1),  
                                train_labels, num_sampled, vocabulary_size))
```

```
    # Optimizer.
```

```
    optimizer = tf.train.AdagradOptimizer(1.0).minimize(loss)
```

```
    # Compute the similarity between minibatch examples and all embeddings.
```

```
    # We use the cosine distance:
```

```
    norm = tf.sqrt(tf.reduce_sum(tf.square(embeddings), 1, keep_dims=True))
```

```
    normalized_embeddings = embeddings / norm
```

```
    valid_embeddings = tf.nn.embedding_lookup(  

```

```

normalized_embeddings, valid_dataset)
similarity = tf.matmul(valid_embeddings, tf.transpose(normalized_embeddings))

num_steps = 100001

with tf.Session(graph=graph) as session:
    tf.initialize_all_variables().run()
    print('Initialized')
    average_loss = 0
    for step in range(num_steps):
        batch_data, batch_labels = generate_batch_cbow(
            batch_size, bag_window)
        feed_dict = {train_dataset : batch_data, train_labels : batch_labels}
        _, l = session.run([optimizer, loss], feed_dict=feed_dict)
        average_loss += l
        if step % 2000 == 0:
            if step > 0:
                average_loss = average_loss / 2000
                # The average loss is an estimate of the loss over the last 2000 batches.
                print('Average loss at step %d: %f' % (step, average_loss))
                average_loss = 0
            # note that this is expensive (~20% slowdown if computed every 500 steps)
            if step % 10000 == 0:
                sim = similarity.eval()
                for i in range(valid_size):
                    valid_word = reverse_dictionary[valid_examples[i]]
                    top_k = 8 # number of nearest neighbors
                    nearest = (-sim[i, :]).argsort()[1:top_k+1]
                    log = 'Nearest to %s:' % valid_word
                    for k in range(top_k):
                        close_word = reverse_dictionary[nearest[k]]
                        log = '%s %s,' % (log, close_word)
                    print(log)
    final_embeddings = normalized_embeddings.eval()

```

Initialized

Average loss at step 0: 8.095695

Nearest to people: ecuador, visibly, contraceptives, bitstream, howell, chagrin,
launcher, circumvented,

Nearest to this: shimura, niki, squealer, tachycardia, sketchy, solubility,
sheer, australis,

Nearest to in: hotly, raions, expended, papen, lachlan, tis, calhoun, stipend,

Nearest to of: mostar, rudolph, divert, berthold, babel, sewage, agi, antony,

.....

Nearest to many: some, several, various, numerous, both, all, hundreds, most,

Nearest to their: its, his, her, your, our, my, the, whose,

Nearest to and: or, but, while, including, however, although, stumble, etc,

Nearest to the: its, his, their, a, any, every, our, each,

In [22]:

```
num_points = 400
```

```
tsne = TSNE(perplexity=30, n_components=2, init='pca', n_iter=5000)
two_d_embeddings = tsne.fit_transform(final_embeddings[1:num_points+1, :])
```

```
def plot(embeddings, labels):  
    assert embeddings.shape[0] >= len(labels), 'More labels than embeddings'  
    pylab.figure(figsize=(15,15)) # in inches  
    for i, label in enumerate(labels):  
        x, y = embeddings[i,:]  
        pylab.scatter(x, y)  
        pylab.annotate(label, xy=(x, y), xytext=(5, 2), textcoords='offset points',  
                        ha='right', va='bottom')  
    pylab.show()
```

```
words = [reverse_dictionary[i] for i in range(1, num_points+1)]
plot(two_d_embeddings, words)
```

```

/home/maxkhk/anaconda/lib/python2.7/site-packages/matplotlib/collections.py:590:
FutureWarning: elementwise comparison failed; returning scalar instead, but in
the future will perform elementwise comparison
    if self. edgecolors == str('face'):

```

